

# Home Assignment 5

## Machine Learning A

FEDERICO FIORIO

October 10, 2022

### 1 Principal Component Analysis (50 points)

#### 1.1 PCA and centering (15 points)

Given what is already proven by the book, where  $Z$  is the centered matrix:

By direct calculation, one can verify that  $Z = X - \mathbf{1}\bar{\mathbf{x}}^T$ . Hence,

$$\bar{\mathbf{z}} = \frac{1}{N}Z^T\mathbf{1} = \frac{1}{N}X^T\mathbf{1} - \frac{1}{N}\bar{\mathbf{x}}\mathbf{1}^T\mathbf{1} = \bar{\mathbf{x}} - \frac{1}{N}\bar{\mathbf{x}} \cdot N = \mathbf{0},$$

Figure 1: Book's proof Logistic regression chapter from Abu-Mostafa et al., Learning from Data (amlbook.com)

From this proof, we can see that the mean vector  $\bar{\mathbf{z}}$  is equal to 0. We can also see that we obtain that as:

$$\frac{1}{N}Z^T\mathbf{1}$$

That operation translates as a sum over the rows of the matrix  $Z$  divided by  $N$ .

This operation it's basically a linear combination of the rows of  $Z$  which is the matrix on which we want to prove that the rank is at most  $d-1$ .

Given what was just said here, we can say that if we sum all the rows of the matrix  $Z$  and divide by  $N$  we get a vector of zeros, which means that at least 1 row in  $Z$  is linearly dependent from the others, which means that the rank of  $Z$  is at most  $d-1$  because there could be at most  $d-1$  linearly independent rows.

## 1.2 Explained variance and Bessel's correction (5 points)

To obtain the eigenvectors and eigenvalues used in PCA, we firstly compute the covariance matrix which as we can see is divided by  $\frac{1}{N}$

$$\mathbf{S} = \frac{1}{N} \sum_{i=1}^N (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^{\top} .$$

Figure 2: covariance matrix from slides

Given the property of eigenvalues:

$cMv = c\lambda v$  where  $c$  is a scalar.

In the above formula,  $c = \frac{1}{N}$  but we can easily change it to  $\frac{1}{N-1}$  and the eigenvalue will also change by that constant  $c$ , but given the explained variance, we can easily see that it won't change even if the constant  $c$  changes:

$$\frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i} = \frac{\sum_{i=1}^k c\lambda_i}{\sum_{i=1}^d c\lambda_i} = \frac{c \sum_{i=1}^k \lambda_i}{c \sum_{i=1}^d \lambda_i} = \frac{\sum_{i=1}^k \lambda_i}{\sum_{i=1}^d \lambda_i}$$

## 1.3 PCA in practice

### 1.3.1 Explained variance (15 points)

Do the PCA:

```
pca = PCA()  
pca.fit(X)  
eigenvalues = pca.singular_values_**2  
explained_var_comp = pca.explained_variance_ratio_
```

```
plt.yscale('log')  
plt.plot(eigenvalues / np.sum(eigenvalues))
```

```
[<matplotlib.lines.Line2D at 0x2428b0e9690>]
```




Figure 3: Plot of eigenvalues / sum of eigenvalues

As we can see from the following image, 10 principal components are not enough to explain 80% of variance

Find out if 10 components are enough to explain 80% of the variance:

```
np.sum(explained_var_comp[:10]) #not enough  
0.7382267688459533
```

```
plt.plot(np.cumsum(explained_var_comp[:10]))  
[<matplotlib.lines.Line2D at 0x242ffd875e0>]
```




Figure 4: Explained variance of 10 principal components




Figure 5: Eigenspectrum, not sure if it's needed

### 1.3.2 Eigendigits (15 points)

```
: for i in range(5):  
    plt.imshow(pca.components_[i].reshape(imshape))  
    plt.show()
```

Figure 6: Code for plotting

Plots:




Figure 7: Eigendigit 1




Figure 8: Eigendigit 2




Figure 9: Eigendigit 3




Figure 10: Eigendigit 4




Figure 11: Eigendigit 5

## 2 Logistic Regression in PyTorch (50 points)

### 2.1 PyTorch self-study (0 points)

## 2.2 Logistic Regression in PyTorch (50 points)

I find it more understandable if I attach directly the screenshots of the code instead of just the lines that I should have written.

```
class LogisticRegressionPytorch(nn.Module):
    def __init__(self, d, m):
        super(LogisticRegressionPytorch, self).__init__()
        # SOMETHING MISSING HERE
        self.linear = nn.Linear(d, m)

    def forward(self, x):
        # SOMETHING MISSING HERE
        outputs = self.linear(x)
        return outputs

logreg_pytorch = LogisticRegressionPytorch(d, m)
print(logreg_pytorch)

LogisticRegressionPytorch(
  (linear): Linear(in_features=2, out_features=4, bias=True)
)
```

Figure 12: Model definition, I did not apply sigmoid layer because crossEntropy already apply that

### Model training and evaluation

First we have to define a loss function. *Double check that output of the network and the expected input of the loss function match!*

```
optimizer = optim.Adam(logreg_pytorch.parameters(), lr=0.01)
# DEFINITION OF LOSS FUNCTION MISSING
CEL = nn.CrossEntropyLoss()
```

Figure 13: Loss function

```

no_epochs = 10000 # Number of training steps
X_train_T = torch.Tensor(X_train) # Automatically casts to float
y_train_T = torch.from_numpy(y_train) # Does not cast to float
#y_train_T = y_train_T.long()
for epoch in range(no_epochs): # Loop over the dataset multiple times
    # Zero the parameter gradients
    optimizer.zero_grad()

    # Forward + backward + optimize
    outputs = logreg_pytorch(X_train_T)
    # SOMETHING MISSING HERE
    loss = CEL(outputs, y_train_T.long()) #probabilities, y_train, if they are long the softmax is done autom.
    # SOMETHING MISSING HERE
    loss.backward()

    optimizer.step()

```

Figure 14: Training loop

```

ws_torch = logreg_pytorch.linear.weight.detach().numpy()
# SOMETHING MISSING HERE
bs_torch = logreg_pytorch.linear.bias.detach().numpy()

```

Figure 15: Getting parameters of the model




Figure 16: visualizing the predictions, it is the same as the sklearn implementation (I know the points are random but it's just to give an idea)